

EM for Linear Regression LAD

AUTHOR

Zehaan Naik

Problem Statement

In the classic linear regression setup, i.e.:

$$y = X\beta + \epsilon \mid \epsilon \sim N(0, 1),$$

we wish to find the LAD estimator for β . For the proof of concept, we consider a simplified version of this problem. Particularly, we seek to minimize:

$$\sum_{i=1}^n |y_i - a - bx_i|$$

Proposed Solution

We know that the following problems have well defined solutions:

$$\sum_{i=1}^n |y_i - a|, \quad \sum_{i=1}^n |y_i - ax_i|$$

Then, one attempt to solve the problem can be looked at in terms breaking the current problem statement into 2 parts:

1. Fix some initial value for a say a^0 and absorb this constant into y_i s and solve the equation:

$$\sum_{i=1}^n |z_i - bx_i| \quad \text{where } z_i = y_i - a^0$$

2. Use the computed value for b say b^1 and absorb this constant ($b^1 x_i$) into y_i s and solve the equation:

$$\sum_{i=1}^n |w_i - a| \quad \text{where } z_i = y_i - b^1 x_i$$

Experiment 1. 2-Variable data set

Generating a random data set

Here, I generate a random 1000 observation data-set with known values for both the target parameters a and b . I also visualize the data on a simple X, y plot.

```
set.seed(123)

### 1. Generate synthetic linear regression dataset
```

```


n <- 1e3 # samples
p <- 2   # features

X <- matrix(rnorm(n * (p-1)), n, (p-1))
X <- cbind(1,X)
beta_true <- c(2,5) # true coefficients
y <- X %*% beta_true + rnorm(n, sd = 1)

## Visualising the data
plot(X[,2], y,
     main = "Generated Data",
     xlab = "X",
     ylab = "y")

```

Generated Data



Sample run of my proposed algorithm to check the parameter fit.

```

### 2. Optimization loop

## Loop Parameters
max_iter <- 100
a <- numeric(max_iter)
b <- numeric(max_iter)
a_curr <- 10
b_curr <- 10

x <- X[,2]

```

```

## Comparison values to store
MAE.loss <- numeric(max_iter)
MAE.params <- numeric(max_iter)

## Optim Loop
for(i in 1:max_iter){
  ## Solving for b_i
  z <- y - a_curr
  z.bar <- z/x
  b_curr <- median(z.bar)
  b[i] <- b_curr

  ## Solving for
  w <- y - b_curr * x
  a_curr <- median(w)
  a[i] <- a_curr

  ## Storing the current MAE on parameters and predictions
  MAE.loss[i] <- mean(abs(y - a_curr - b_curr*x))
  MAE.params[i] <- mean(abs(c(a_curr, b_curr) - beta_true))
}

cat("True parameters are a = ", beta_true[1] , "and b = ", beta_true[2])

```

True parameters are a = 2 and b = 5

```
cat("Post optimization we get a = ", a_curr , "and b = ", b_curr)
```

Post optimization we get a = 2.039266 and b = 5.120642

Visualizing the results

Here, I plot the MAE on parameters, prediction losses and the actual model fit to check how well the methodology performs on our generated data set. We can observe that we have a quick convergence of both the loss values and the MAE on parameters indicating a stable fit by the algorithm.


```

#### 3. Plots to visualise method performance

## Model Fit
plot(x,y,
     main = "EM for LAD (model fit)",
     xlab = "X",
     ylab = "y",
     pch = 19)
lines(x, a_curr + b_curr*x, col = "red")
legend("bottomright", c("Line of best fit"), col = ("red"), lwd = 2)


```

EM for LAD (model fit)




```
## Checking the MAE over predictions
plot(1:max_iter, MAE.loss[1:max_iter], col = "red", type = "l",
     main = "MAE on predictions across iterations of model fitting",
     xlab = "Iters",
     ylab = "MAE")
```

MAE on predictions across iterations of model fitting



```
## Checking the MAE loss over params
plot(1:max_iter, MAE.params[1:max_iter], col = "blue", type = "l",
     main = "MAE on parameters across iterations of model fitting",
     xlab = "Iters",
     ylab = "MAE")
```

MAE on parameters across iterations of model fitting



Experiment 1. p -Covariate data set

Now, just to complicate the problem a little bit, I generate a data with p many dimensions in X and try to estimate all the p parameters using this new methodology. The overall approach remains the same, we fix all other parameters while trying to optimize for an isolated parameter.

```
set.seed(123)

### 1. Generate synthetic linear regression dataset
n <- 1e3 # samples
p <- 10  # features

X <- matrix(rnorm(n * (p-1)), n, (p-1))
X <- cbind(1, X)
beta_true <- c(2, 1:(p-1)) # true coefficients
y <- X %*% beta_true + rnorm(n, sd = 1)

### 2. Optimization loop

## Loop Parameters
max_iter <- 100
a <- numeric(max_iter)
b <- matrix(0, nrow = max_iter, ncol = p-1)
a_curr <- 1
b_curr <- rep(1, p-1)
```

```

## Comparision values to store
MAE.loss <- numeric(max_iter)
MAE.params <- numeric(max_iter)

## Optim Loop
for(i in 1:max_iter){
  for(j in 1:(p-1)){
    x.j <- X[,1+j]
    X.rest <- X[,-c(1,1+j)]
    ## Solving for b_i,j
    z <- y - a_curr - X.rest %*% b_curr[-j]
    z.bar <- z/x.j
    b_curr[j] <- median(z.bar)
    b[i,] <- b_curr
  }

  ## Solving for a_i
  w <- y - X[, -1] %*% b_curr
  a_curr <- median(w)
  a[i] <- a_curr

  ## Storing the current MAE on parameters and predictions
  MAE.loss[i] <- mean(abs(y - a_curr - X[, -1] %*% b_curr))
  MAE.params[i] <- mean(abs(c(a_curr, b_curr) - beta_true))
}
### CHECKING THE VALUES OF ALL CALCULATED PARAMETERS
print("True Parameters:")

```

```
[1] "True Parameters:"
```

```
print(beta_true)
```

```
[1] 2 1 2 3 4 5 6 7 8 9
```

```
print("Estimated Parameters:")
```

```
[1] "Estimated Parameters:"
```

```
print(c(a_curr, b_curr))
```


```
[1] 1.936893 1.004593 2.043029 2.954753 4.078318 4.970971 5.981237 7.002671
[9] 7.967983 9.008431
```

```

### 3. Plotting the results
## Checking the MAE over predictions
plot(1:max_iter, MAE.loss[1:max_iter], col = "red", type = "l",
     main = "MAE on predictions across iterations of model fitting",
     xlab = "Iters",
     ylab = "MAE")


```

MAE on predictions across iterations of model fitting



```
## Checking the MAE loss over params
plot(1:max_iter, MAE.params[1:max_iter], col = "blue", type = "l",
     main = "MAE on parameters across iterations of model fitting",
     xlab = "Iters",
     ylab = "MAE")
```


MAE on parameters across iterations of model fitting



Clearly, we are able to compute all the parameters with reasonable accuracy and the convergence of MAEs for both parameters and prediction losses indicate a stable fit.

Conclusion

This is a promising method for estimating the LAD estimates for parameters. Given that we are able to work out the underlying theory to prove the descent condition for the algorithm set-up, this method can provide a reliable framework to optimize regression problems in the absolute deviation framework which often lacks an analytical solution.